

### **Ex No: 3 BUILD A CONVOLUTIONAL NEURAL NETWORK**

Aim:

To build a simple convolutional neural network with Keras/TensorFlow.

Procedure:

1. Download and load the dataset.
2. Perform analysis and preprocessing of the dataset.
3. Build a simple neural network model using Keras/TensorFlow.
4. Compile and fit the model.
5. Perform prediction with the test dataset.
6. Calculate performance metrics.

Program:

```
import tensorflow as tf

from tensorflow.keras import datasets, layers, models

import matplotlib.pyplot as plt

(train_images, train_labels), (test_images, test_labels) = datasets.cifar10.load_data()

train_images, test_images = train_images / 255.0, test_images / 255.0

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

plt.figure(figsize=(10,10))

for i in range(25):

    plt.subplot(5,5,i+1)

    plt.xticks([])

    plt.yticks([])

    plt.grid(False)

    plt.imshow(train_images[i])

    plt.xlabel(class_names[train_labels[i][0]])

model = models.Sequential()

model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
```

```
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.Flatten())
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dense(10))
model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])
history = model.fit(train_images, train_labels, epochs=10,
                  validation_data=(test_images, test_labels))
plt.plot(history.history['accuracy'], label='accuracy')
plt.plot(history.history['val_accuracy'], label = 'val_accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.ylim([0.5, 1])
plt.legend(loc='lower right')
test_loss, test_acc = model.evaluate(test_images, test_labels, verbose=2)
print(test_acc)
```

Output:

```
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 — 2s 0us/step
/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `in
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/10
1563/1563 — 76s 47ms/step - accuracy: 0.3427 - loss: 1.7721 - val_accuracy: 0.5313 - val_loss: 1.2924
Epoch 2/10
1563/1563 — 71s 45ms/step - accuracy: 0.5636 - loss: 1.2230 - val_accuracy: 0.6222 - val_loss: 1.0709
Epoch 3/10
1563/1563 — 85s 47ms/step - accuracy: 0.6342 - loss: 1.0308 - val_accuracy: 0.6458 - val_loss: 1.0148
Epoch 4/10
1563/1563 — 78s 45ms/step - accuracy: 0.6789 - loss: 0.9147 - val_accuracy: 0.6804 - val_loss: 0.9242
Epoch 5/10
1563/1563 — 82s 45ms/step - accuracy: 0.7047 - loss: 0.8401 - val_accuracy: 0.6913 - val_loss: 0.8974
Epoch 6/10
1563/1563 — 83s 46ms/step - accuracy: 0.7254 - loss: 0.7828 - val_accuracy: 0.6806 - val_loss: 0.9246
Epoch 7/10
1563/1563 — 81s 45ms/step - accuracy: 0.7442 - loss: 0.7281 - val_accuracy: 0.6870 - val_loss: 0.9263
Epoch 8/10
1563/1563 — 72s 46ms/step - accuracy: 0.7592 - loss: 0.6890 - val_accuracy: 0.6826 - val_loss: 0.9506
Epoch 9/10
1563/1563 — 80s 45ms/step - accuracy: 0.7729 - loss: 0.6450 - val_accuracy: 0.6994 - val_loss: 0.8676
Epoch 10/10
1563/1563 — 73s 47ms/step - accuracy: 0.7872 - loss: 0.6034 - val_accuracy: 0.7073 - val_loss: 0.8830
313/313 - 3s - 11ms/step - accuracy: 0.7073 - loss: 0.8830
0.7073000073432922
```

## Result:

Thus the program for building a simple convolutional neural network was executed successfully.